

Isolated inference latency (batch = 96, averaged over 6 tasks)

